

## COMS 573 Homework 2

---

Boudhayan Banerjee (715909222)

27 March, 2016

1. a. *set.seed*(1)  
     $x1 = \text{runif}(100)$   
     $x2 = 0.5 * x1 + \text{rnorm}(100)/10$   
     $y = 2 + 2 * x1 + 0.3 * x2 + \text{rnorm}(100)$

Form of the linear model,

$$Y = 2 + 2X_1 + 0.3X_2 + \epsilon$$

Regression coefficients,


$$\beta_0 = 2, \beta_1 = 2, \beta_2 = 0.3$$

- b. To find correlation between  $x1$  and  $x2$  we run the following command:  
     $\text{cor}(x1, x2)$

The correlation is, 0.8351212.

**Scatter-plot :**

We display the scatter plot we run the command,  $\text{plot}(x1, x2)$ .



- c. To fit a least square regression to predict Y using x1 and x2 we need to use following command:

```
lm.fit = lm(y ~ x1 + x2)
summary(lm.fit)
```

Results:

---

Call:

```
lm(formula = y ~ x1 + x2)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.75207	-0.59159	-0.03182	0.64792	2.26200

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.4688	0.2121	11.641	<2e-16 ***
x1	1.3295	0.6596	2.016	0.0466 *
x2	0.2163	1.0369	0.209	0.8352

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.966 on 97 degrees of freedom

Multiple R-squared: 0.1408, Adjusted R-squared: 0.1231

F-statistic: 7.947 on 2 and 97 DF, p-value: 0.0006365

Regression coefficients are,

$$\hat{\beta}_0 = 2.4688, \hat{\beta}_1 = 1.3295, \hat{\beta}_2 = 0.2163$$

The regression coefficients are relatively close to the true coefficients. But the standard error of this model is high.

We can reject null hypothesis  $H_0 : \beta_1 = 0$ , as the p value is 0.0466. Which is much less than 5%.

We can not reject null hypothesis  $H_0 : \beta_2 = 0$ . Because the p value is 0.8352 which is more than standard cut-off 5%.

- d. To fit a least square regression to predict Y using only x1 we need to use following command:

```
lm.fit = lm(y ~ x1)
summary(lm.fit)
```

Results:

```
Call:
lm(formula = y ~ x1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.75787	-0.56934	-0.04168	0.64766	2.24483

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.4649	0.2102	11.725	< 2e-16 ***
x1	1.4445	0.3610	4.001	0.000123 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9612 on 98 degrees of freedom

Multiple R-squared: 0.1404, Adjusted R-squared: 0.1316

F-statistic: 16.01 on 1 and 98 DF, p-value: 0.0001226

Yes, we can reject the null hypothesis for the regression coefficient as the p-value 0.000123 for its t-statistic is very small.

- e. To fit a least square regression to predict Y using only x2 we need to use following command:

```
lm.fit = lm(y ~ x2)
```

```
summary(lm.fit)
```

Results:

Call:

```
lm(formula = y ~ x2)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.90820	-0.63486	0.05152	0.62718	2.24278

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.7084	0.1783	15.186	< 2e-16 ***
x2	1.9618	0.5792	3.387	0.00102 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9809 on 98 degrees of freedom

Multiple R-squared: 0.1048, Adjusted R-squared: 0.09566

F-statistic: 11.47 on 1 and 98 DF, p-value: 0.001018

Yes, we can reject the null hypothesis for the regression coefficient as the p-value 0.00102 for its t-statistic is very small.

- f. The results obtained in part c,d,e does not contradict each other. From part b we can see that x1 and x2 are collinear. Therefore when we try to fit the linear regression model with both x1 and x2 it is difficult to spot individual effect. But when we fit the same model separately on x1 and x2 then the relationship between Y and each predictor is more visible.

- g. We modify our model to accommodate another observation. We run the following commands.

```
x1 = c(x1,0.1)
```

```
x2 = c(x2,0.8)
```

```
y = c(y,6)
```

```
lm.fitm = lm(y ~ x1 + x2)
```

```
summary(lm.fitm)
```

Result:

```
Call:
lm(formula = y ~ x1 + x2)

Residuals:
    Min       1Q   Median       3Q      Max
-2.7825 -0.6509  0.0189  0.6767  2.3316

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   2.5751      0.2137  12.048  <2e-16 ***
x1             0.3346      0.5471   0.612  0.5422
x2             1.8798      0.8293   2.267  0.0256 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9928 on 98 degrees of freedom
Multiple R-squared:  0.1454,    Adjusted R-squared:  0.128
F-statistic: 8.336 on 2 and 98 DF,  p-value: 0.0004535
```

```
lm.fitmx1 = lm(y ~ x1)
summary(lm.fitmx1)
```

Result:

```
Call:
lm(formula = y ~ x1)

Residuals:
    Min       1Q   Median       3Q      Max
-2.8881 -0.5963 -0.0340  0.6105  3.2772

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   2.5977      0.2179  11.921  < 2e-16 ***
x1             1.2513      0.3760   3.327  0.00123 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.013 on 99 degrees of freedom
Multiple R-squared:  0.1006,    Adjusted R-squared:  0.0915
F-statistic: 11.07 on 1 and 99 DF,  p-value: 0.001231
```

```
lm.fitmx2 = lm(y ~ x2)
summary(lm.fitmx2)
```

Result:

```
Call:
lm(formula = y ~ x2)

Residuals:
    Min       1Q   Median       3Q      Max
-2.83733 -0.62236  0.04366  0.61993  2.31160

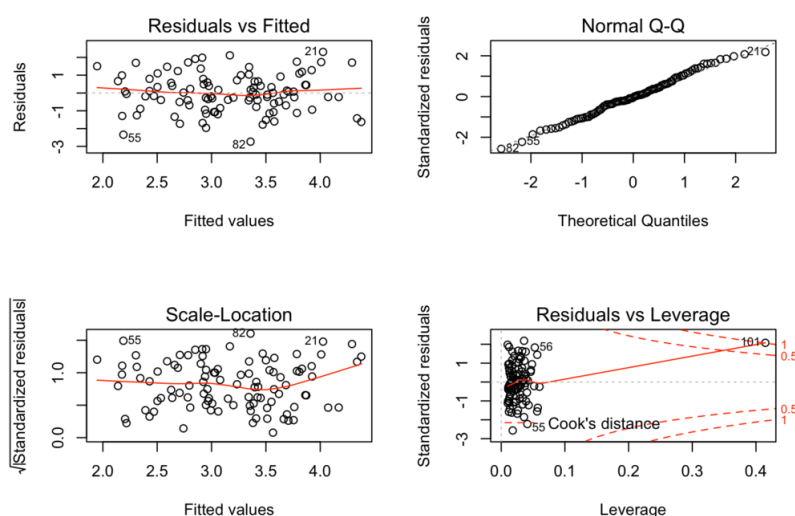
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   2.6486     0.1762   15.03 < 2e-16 ***
x2             2.2547     0.5567    4.05 0.000102 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9897 on 99 degrees of freedom
Multiple R-squared:  0.1421,    Adjusted R-squared:  0.1335
F-statistic: 16.4 on 1 and 99 DF,  p-value: 0.0001019
```

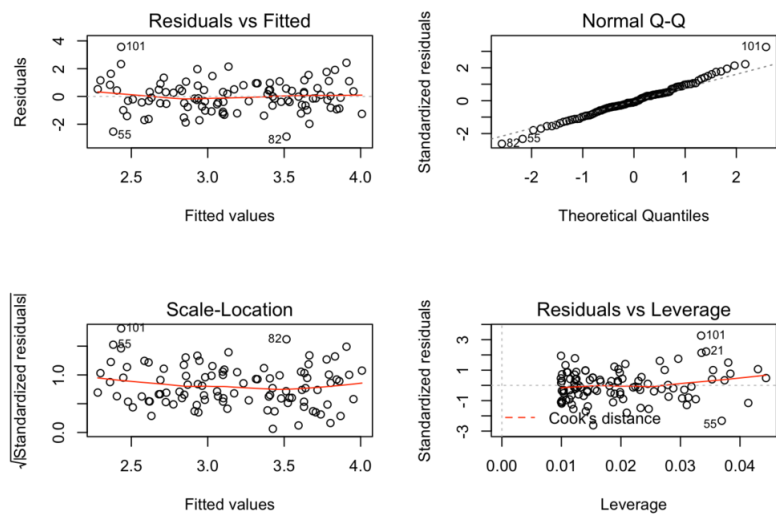
OBSERVATION :

*In the first model, x1 becomes statistically insignificant and x2 becomes statistically significant because of the change in p-values between the two linear regressions.*

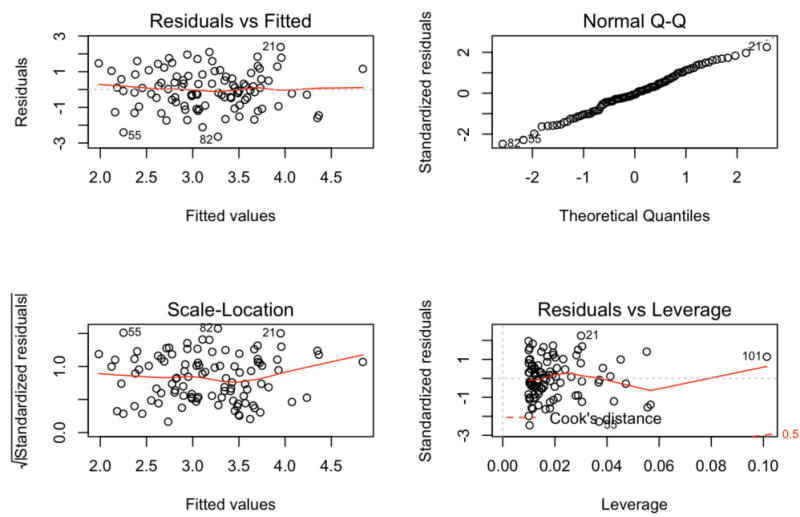
```
par(mfrow=c(2,2))
plot(lm.fitm)
```




```
par(mfrow=c(2,2))
plot(lm.fitmx1)
```




```
par(mfrow=c(2,2))
plot(lm.fitmx2)
```




```
plot(predict(lm.fitm), rstudent(lm.fitm))
```



`plot(predict(lm.fitmx1), rstudent(lm.fitmx1))`



`plot(predict(lm.fitmx2), rstudent(lm.fitmx2))`



In the model where we use both x1 and x2 as predictor the last point is high lever-



age point.

In the model where we use only x1 as predictor the last point is an outlier.

In the model where we use only x2 as predictor the last point is a high leverage point.

2. To use Bayes classifier, we have to find the class (k) for which,

$$p_k(x) = \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma_k} \exp(-\frac{1}{2\sigma_k^2}(x - \mu_k)^2)}{\sum \pi_l \frac{1}{\sqrt{2\pi}\sigma_l} \exp(-\frac{1}{2\sigma_l^2}(x - \mu_l)^2)}$$

is largest. As the log function is monotonically increasing, it is equivalent to finding k for which,

$$\log(p_k(x)) = \frac{\log(\pi_k) + \log(\frac{1}{\sqrt{2\pi}\sigma_k}) + -\frac{1}{2\sigma_k^2}(x - \mu_k)^2}{\log(\sum \pi_l \frac{1}{\sqrt{2\pi}\sigma_l} \exp(-\frac{1}{2\sigma_l^2}(x - \mu_l)^2))}$$

is largest. Thus we have,

$$\log(p_k(x)) \log(\sum \pi_l \frac{1}{\sqrt{2\pi}\sigma_l} \exp(-\frac{1}{2\sigma_l^2}(x - \mu_l)^2)) = \log(\pi_k) + \log(\frac{1}{\sqrt{2\pi}\sigma_k}) + -\frac{1}{2\sigma_k^2}(x - \mu_k)^2$$

Therefore,

$$\delta(x) = \log(\pi_k) + \log(\frac{1}{\sqrt{2\pi}\sigma_k}) + -\frac{1}{2\sigma_k^2}(x - \mu_k)^2$$

As we can see that  $\delta(x)$  is a quadratic function of x. Therefore the Bayes classifier is not linear but quadratic.

3. We have,

$$p_k(x) = \frac{\pi_k \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}(x - \mu_k)^2)}{\sum \pi_l \frac{1}{\sqrt{2\pi}\sigma} \exp(-\frac{1}{2\sigma^2}(x - \mu_l)^2)}$$

Now class k contains both yes and no. Therefore we have,

$$p_{yes}(x) = \frac{\pi_{yes} \exp(-\frac{1}{2\sigma^2}(x - \mu_{yes})^2)}{\sum \pi_l \exp(-\frac{1}{2\sigma^2}(x - \mu_l)^2)}$$

It is given that 80% companies issued dividend, hence  $\pi_{yes} = 0.80$  and  $\pi_{no} = 0.20$

It is also given that  $\sigma^2 = 36$

Also mean value of X for companies that issued dividend is  $\mu_{yes} = 10$   
Therefore substituting values in the original equation we get,

$$= \frac{\pi_{yes} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_{yes})^2\right)}{\pi_{yes} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_{yes})^2\right) + \pi_{no} \exp\left(-\frac{1}{2\sigma^2}(x - \mu_{no})^2\right)}$$

$$= \frac{0.80 \exp\left(-\frac{1}{2*36}(x - 10)^2\right)}{0.80 \exp\left(-\frac{1}{2*36}(x - 10)^2\right) + 0.20 \exp\left(-\frac{1}{2*36}x^2\right)}$$

For x=4 we have,

$$p_{yes}(4) = \frac{0.80 \exp\left(-\frac{1}{2*36}(4 - 10)^2\right)}{0.80 \exp\left(-\frac{1}{2*36}(4 - 10)^2\right) + 0.20 \exp\left(-\frac{1}{2*36}4^2\right)} = 75.2\%$$

4. a. First we produce the summary of the weekly data set,

```
library(ISLR)
summary(Weekly)
```

Year	Lag1	Lag2	Lag3
Min. :1990	Min. :-18.195	Min. :-18.195	Min. :-18.195
1st Qu.:1995	1st Qu.: -1.154	1st Qu.: -1.154	1st Qu.: -1.158
Median :2000	Median : 0.241	Median : 0.241	Median : 0.241
Mean :2000	Mean : 0.151	Mean : 0.151	Mean : 0.147
3rd Qu.:2005	3rd Qu.: 1.405	3rd Qu.: 1.409	3rd Qu.: 1.409
Max. :2010	Max. : 12.026	Max. : 12.026	Max. : 12.026

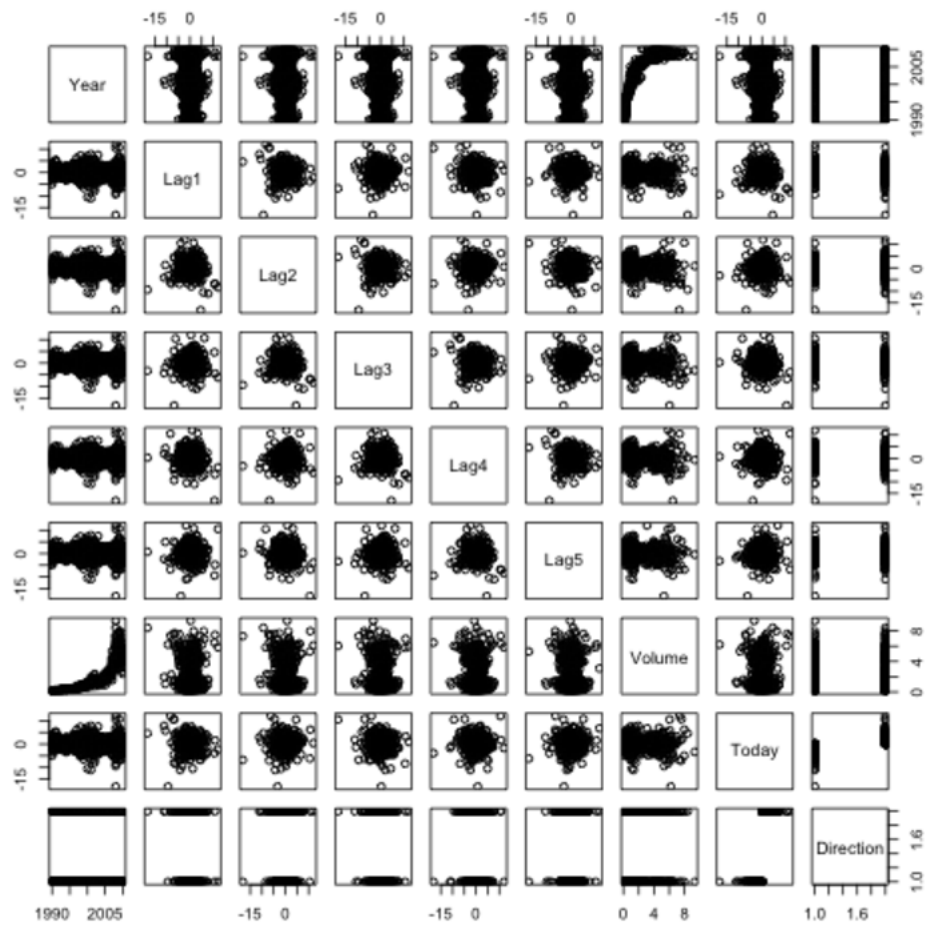
  

Lag4	Lag5	Volume	Today
Min. :-18.195	Min. :-18.195	Min. :0.087	Min. :-18.195
1st Qu.: -1.158	1st Qu.: -1.166	1st Qu.:0.332	1st Qu.: -1.154
Median : 0.238	Median : 0.234	Median :1.003	Median : 0.241
Mean : 0.146	Mean : 0.140	Mean :1.575	Mean : 0.150
3rd Qu.: 1.409	3rd Qu.: 1.405	3rd Qu.:2.054	3rd Qu.: 1.405
Max. : 12.026	Max. : 12.026	Max. :9.328	Max. : 12.026

Direction  
Down:484  
Up :605

Next we produce the graphical representations,

```
pairs(Weekly)
```



```
cor(Weekly[, -9])
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume
Year	1.00000	-0.032289	-0.03339	-0.03001	-0.031128	-0.030519	0.84194
Lag1	-0.03229	1.000000	-0.07485	0.05864	-0.071274	-0.008183	-0.06495
Lag2	-0.03339	-0.074853	1.00000	-0.07572	0.058382	-0.072499	-0.08551
Lag3	-0.03001	0.058636	-0.07572	1.00000	-0.075396	0.060657	-0.06929
Lag4	-0.03113	-0.071274	0.05838	-0.07540	1.000000	-0.075675	-0.06107
Lag5	-0.03052	-0.008183	-0.07250	0.06066	-0.075675	1.000000	-0.05852
Volume	0.84194	-0.064951	-0.08551	-0.06929	-0.061075	-0.058517	1.00000
Today	-0.03246	-0.075032	0.05917	-0.07124	-0.007826	0.011013	-0.03308
Today							
Year	-0.032460						
Lag1	-0.075032						
Lag2	0.059167						
Lag3	-0.071244						
Lag4	-0.007826						
Lag5	0.011013						
Volume	-0.033078						
Today	1.000000						

It is evident from the results that Year and Volume are correlated.

b. `attach(Weekly)`

```
glm.fit = glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = Weekly,
family = binomial)
summary(glm.fit)
```

```
Call:
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
    Volume, family = binomial, data = Weekly)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.695   -1.256    0.991    1.085    1.458

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)   0.2669    0.0859   3.11  0.0019 **
Lag1          -0.0413    0.0264  -1.56  0.1181
Lag2           0.0584    0.0269   2.18  0.0296 *
Lag3          -0.0161    0.0267  -0.60  0.5469
Lag4          -0.0278    0.0265  -1.05  0.2937
Lag5          -0.0145    0.0264  -0.55  0.5833
Volume        -0.0227    0.0369  -0.62  0.5377
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1496.2  on 1088  degrees of freedom
Residual deviance: 1486.4  on 1082  degrees of freedom
AIC: 1500

Number of Fisher Scoring iterations: 4
```

Lag2 is the predictor that has some statistical significance as its p-value is 0.0296, which is less than 5% cut off.

```
c. glm.probs = predict(glm.fit, type = "response")
glm.pred = rep("Down", length(glm.probs))
glm.pred[glm.probs > 0.5] = "Up"
table(glm.pred, Direction)
```

```
      Direction
glm.pred Down  Up
Down     54   48
Up      430  557
```

Percentage of correct predictions on the training data is  $(54 + 557) / (54 + 557 + 48 + 430) = 56.106\%$ .

For those weeks when market goes up the percentage of correct predictions is  $557 / (557 + 48) = 92.066\%$

.

For those weeks when market goes down the percentage of correct predictions is  $54 / (430 + 54) = 11.157\%$ .

```
d. train <- (Year < 2009)
Weekly.20092010 <- Weekly[!train, ]
Direction.20092010 <- Direction[!train]
fit.glm2 <- glm(Direction ~ Lag2, data = Weekly, family=binomial, subset = train)
summary(fit.glm2)
```

```
Call:
glm(formula = Direction ~ Lag2, family = binomial, data = Weekly,
     subset = train)
```

```
Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.536  -1.264   1.021   1.091   1.368
```

```
Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  0.20326    0.06428   3.162  0.00157 **
Lag2         0.05810    0.02870   2.024  0.04298 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(Dispersion parameter for binomial family taken to be 1)
```

```
Null deviance: 1354.7 on 984 degrees of freedom
Residual deviance: 1350.5 on 983 degrees of freedom
AIC: 1354.5
```

```
Number of Fisher Scoring iterations: 4
```

```
Now, probs2 <- predict(fit.glm2, Weekly.20092010, type = "response")
pred.glm2 <- rep("Down", length(probs2))
pred.glm2[probs2 > 0.5] <- "Up"
table(pred.glm2, Direction.20092010)
```

```
              Direction.20092010
pred.glm2 Down Up
      Down     9  5
      Up     34 56
```

Percentage of correct prediction on training data =  $(9 + 56) / 104 = 62.5\%$ .

Correct prediction percentage:

For week when the market goes up =  $56 / (56 + 5) = 91.80\%$

For week when the market goes down =  $9 / (9 + 34) = 20.93\%$

e. `library(MASS)`  
`lda.fit = lda(Direction ~ Lag2, data = Weekly, subset = train)`  
`lda.pred = predict(lda.fit, Weekly.0910)`  
`table(lda.pred$class, Direction.0910)`

	Direction.0910	
	Down	Up
Down	9	5
Up	34	56

`mean(lda.pred$class == Direction.0910)`

we get mean = 0.625. Therefore percentage of correct prediction on training data is 62.5%.

Correct prediction percentage:

For week when the market goes up =  $56 / (56 + 5) = 91.80\%$

For week when the market goes down =  $9 / (9 + 34) = 20.93\%$

f. `qda.fit = qda(Direction ~ Lag2, data = Weekly, subset = train)`  
`qda.class = predict(qda.fit, Weekly.0910)$class`  
`table(qda.class, Direction.0910)`

	Direction.0910	
	Down	Up
qda.class Down	0	0
Up	43	61

`mean(qda.class == Direction.0910)`

we get mean = 0.5865. Therefore percentage of correct prediction on training data is 58.65%.

Correct prediction percentage:

For week when the market goes up =  $61 / (61) = 100\%$

For week when the market goes down =  $0 / (43) = 0\%$

g. In case of both QDA and LDA we have the assumption that all the classes come from a multivariate normal. There are couple of things we need to check before using QDA.

a. We should ensure that the class means are significantly different.

Here we have single class, hence we should have no problem using QDA.

b. We should also check if we have sufficient number of data points in the training data set.

In this case we have 985 data points. Therefore we can use QDA.

To ensure that all the classes come from a multivariate normal we should perform `HZ.test()` before using QDA.

```
h. library(class)
train.X = as.matrix(Lag2[train])
test.X = as.matrix(Lag2[!train])
train.Direction = Direction[train]
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 1)
table(knn.pred, Direction.0910)
```

```
mean(knn.pred == Direction.0910)
```

```
      Direction.0910
knn.pred Down Up
Down      21 30
Up        22 31
```

we get mean = 0.5. Therefore percentage of correct prediction on training data is 50%.

Correct prediction percentage:

For week when the market goes up =  $31 / (61) = 50.82\%$

For week when the market goes down =  $21 / (43) = 48.84\%$

- i. Logistic regression and LDA both provides best result on this data with correct prediction percentage of 62.5% i.e 37.5% error rate.
- j. Naive Bayes can be considered a better classifier due to it's robustness. But whether it can actually deliver better result depends on various other factors.